

California_housing

August 14, 2023

```
[ ]: from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout
from tensorflow.keras.wrappers.scikit_learn import KerasRegressor, KerasClassifier
from tensorflow.keras.callbacks import EarlyStopping, ModelCheckpoint

import matplotlib.pyplot as plt
import seaborn as sns

import numpy as np
import pandas as pd

from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.datasets import fetch_california_housing
from sklearn.preprocessing import MinMaxScaler
from sklearn.metrics import precision_score, f1_score, accuracy_score, recall_score, confusion_matrix
```

1 Regression problem

1.1 Loading dataset

Load the Boston Housing dataset from the Keras library.

Due to the ethical issues surrounding the Boston housing dataset, primarily the inclusion of a racist feature engineered by the dataset's authors, the Boston housing dataset will not be used. Alternatively, the California housing dataset, available in SKlearn, will be used instead.

```
[ ]: ca_housing = fetch_california_housing(as_frame=True)
```

```
[ ]: feature_df = ca_housing.data
target_df = ca_housing.target
```

1.2 Explore and preprocess

Explore and preprocess the data (e.g., normalization, one-hot encoding, etc.).

```
[ ]: (num_samples, num_features) = feature_df.shape
```

```
print(f'Sample size: {num_samples}')
print(f'Number of features: {num_features}')
```

Sample size: 20640

Number of features: 8

```
[ ]: # concatenate all train and test data and targets together
full_df = pd.concat([feature_df, target_df], axis=1)
```

```
[ ]: full_df
```

```
[ ]:
      MedInc  HouseAge  AveRooms  AveBedrms  Population  AveOccup  Latitude  \
0      8.3252      41.0  6.984127   1.023810      322.0    2.555556      37.88
1      8.3014      21.0  6.238137   0.971880     2401.0    2.109842      37.86
2      7.2574      52.0  8.288136   1.073446      496.0    2.802260      37.85
3      5.6431      52.0  5.817352   1.073059      558.0    2.547945      37.85
4      3.8462      52.0  6.281853   1.081081      565.0    2.181467      37.85
...
20635  1.5603      25.0  5.045455   1.133333      845.0    2.560606      39.48
20636  2.5568      18.0  6.114035   1.315789      356.0    3.122807      39.49
20637  1.7000      17.0  5.205543   1.120092     1007.0    2.325635      39.43
20638  1.8672      18.0  5.329513   1.171920      741.0    2.123209      39.43
20639  2.3886      16.0  5.254717   1.162264     1387.0    2.616981      39.37

      Longitude  MedHouseVal
0      -122.23         4.526
1      -122.22         3.585
2      -122.24         3.521
3      -122.25         3.413
4      -122.25         3.422
...
20635    -121.09         0.781
20636    -121.21         0.771
20637    -121.22         0.923
20638    -121.32         0.847
20639    -121.24         0.894
```

[20640 rows x 9 columns]

1.2.1 Correlation Matrix


```
[ ]: correlation_matrix = full_df.corr()

corr_w_target = correlation_matrix['MedHouseVal']
```

```
[ ]: plt.figure(figsize=(12,10))

sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', center=0)

plt.title('Correlation Matrix Heatmap')
plt.show()
```



```
[ ]: print(f'Correlations with the target variable:\n{corr_w_target}')
```

Correlations with the target variable:

```
MedInc      0.688075
HouseAge    0.105623
AveRooms     0.151948
AveBedrms   -0.046701
Population  -0.024650
AveOccup    -0.023737
```

```
Latitude      -0.144160
Longitude     -0.045967
MedHouseVal    1.000000
Name: MedHouseVal, dtype: float64
```

1.2.2 Feature histplots

```
[ ]: num_rows = 4
      num_cols = 2

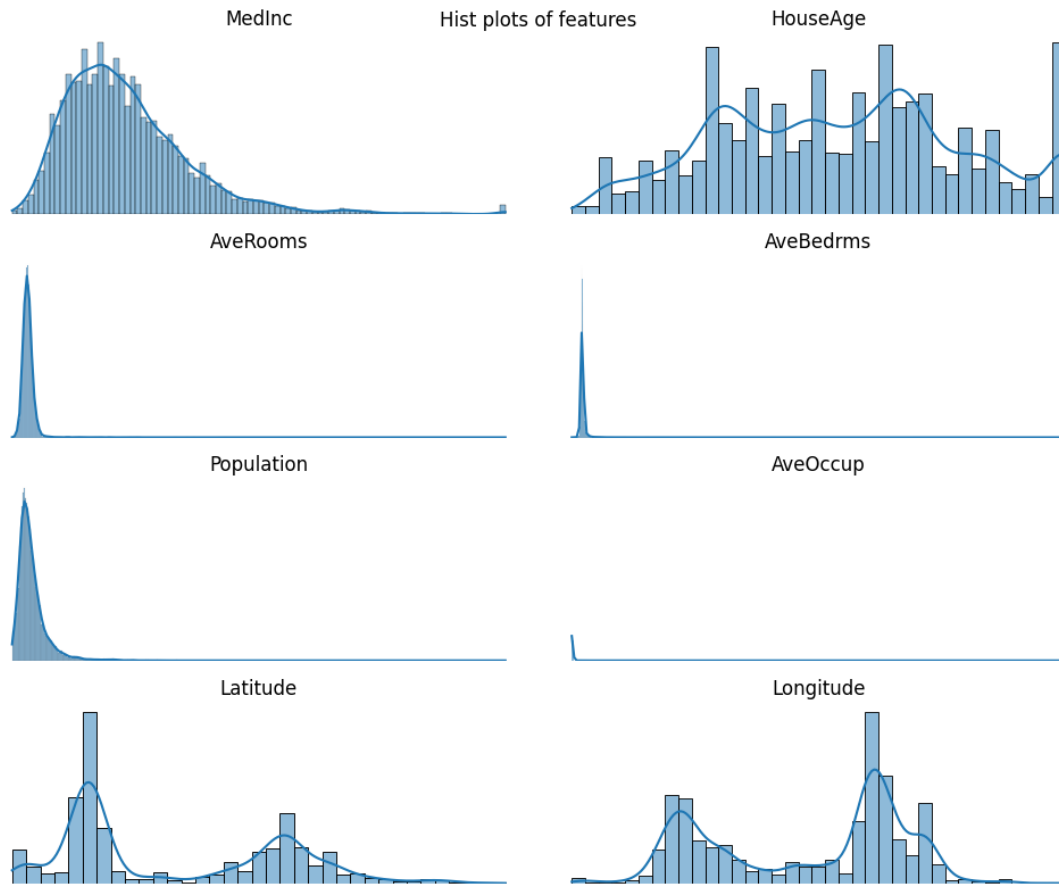
      fig, axes = plt.subplots(num_rows, num_cols, figsize=(10,8))
      fig.subplots_adjust(hspace=0.5)

      for i in range(num_features):
          row_idx = i // num_cols
          col_idx = i % num_cols

          sns.histplot(feature_df.iloc[:,i], ax=axes[row_idx, col_idx], kde=True)
          axes[row_idx, col_idx].set_title(f'{feature_df.columns[i]}')

          if i >= num_features - (num_rows * num_cols):
              axes[row_idx, col_idx].axis('off')

      plt.tight_layout()
      plt.suptitle('Hist plots of features')
      plt.show()
```



1.2.3 Feature boxplots

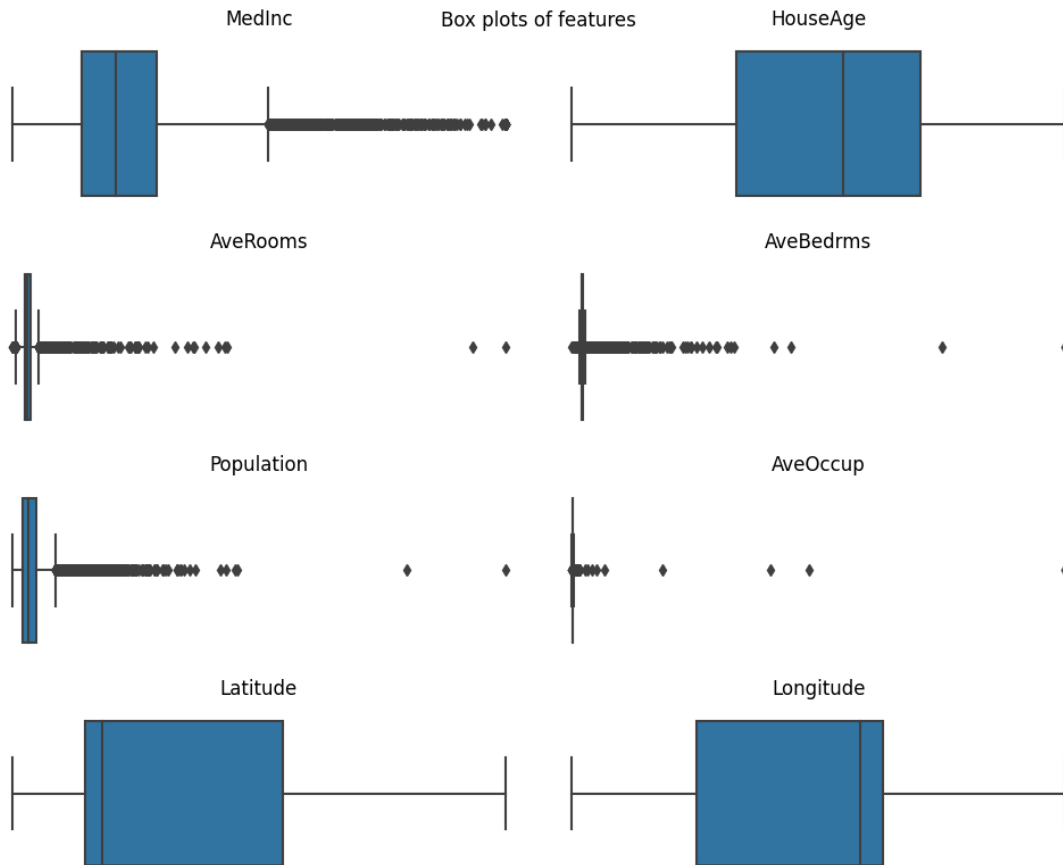
```
[ ]: fig, axes = plt.subplots(num_rows, num_cols, figsize=(10,8))
fig.subplots_adjust(hspace=0.5)

for i in range(num_features):
    row_idx = i // num_cols
    col_idx = i % num_cols

    sns.boxplot(x=feature_df.iloc[:,i], ax=axes[row_idx, col_idx])
    axes[row_idx, col_idx].set_title(f'{feature_df.columns[i]}')

    if i >= num_features - (num_rows * num_cols):
        axes[row_idx, col_idx].axis('off')

plt.tight_layout()
plt.suptitle('Box plots of features')
plt.show()
```



1.2.4 Address significant outliers

```
[ ]: full_df.describe()
```

```
[ ]:
count    MedInc    HouseAge    AveRooms    AveBedrms    Population  \
count    20640.000000    20640.000000    20640.000000    20640.000000    20640.000000
mean      3.870671      28.639486      5.429000      1.096675     1425.476744
std       1.899822      12.585558      2.474173      0.473911     1132.462122
min       0.499900       1.000000      0.846154      0.333333       3.000000
25%       2.563400      18.000000      4.440716      1.006079      787.000000
50%       3.534800      29.000000      5.229129      1.048780     1166.000000
75%       4.743250      37.000000      6.052381      1.099526     1725.000000
max      15.000100      52.000000     141.909091     34.066667    35682.000000

count    AveOccup    Latitude    Longitude    MedHouseVal
count    20640.000000    20640.000000    20640.000000    20640.000000
mean      3.070655      35.631861     -119.569704      2.068558
std      10.386050       2.135952       2.003532       1.153956
min       0.692308      32.540000     -124.350000       0.149990
```

25%	2.429741	33.930000	-121.800000	1.196000
50%	2.818116	34.260000	-118.490000	1.797000
75%	3.282261	37.710000	-118.010000	2.647250
max	1243.333333	41.950000	-114.310000	5.000010

```
[ ]: full_df[full_df['AveRooms'] > 50]
```

```
[ ]:
      MedInc  HouseAge  AveRooms  AveBedrms  Population  AveOccup  \
1912  4.9750      16.0  56.269231  10.153846      54.0  2.076923
1913  4.0714      19.0  61.812500  11.000000     112.0  2.333333
1914  1.8750      33.0 141.909091  25.636364      30.0  2.727273
1979  4.6250      34.0 132.533333  34.066667      36.0  2.400000
2395  3.8750      23.0  50.837838  10.270270      64.0  1.729730
9676  3.2431      14.0  52.848214  11.410714     265.0  2.366071
11707 1.1912      22.0  52.690476   8.857143      98.0  2.333333
11862 2.6250      25.0  59.875000  15.312500      28.0  1.750000
12447 1.6154      17.0  62.422222  14.111111      83.0  1.844444
```

	Latitude	Longitude	MedHouseVal
1912	39.01	-120.16	2.06300
1913	39.01	-120.06	4.37500
1914	38.91	-120.10	5.00001
1979	38.80	-120.08	1.62500
2395	37.12	-119.34	1.25000
9676	37.64	-119.02	2.21400
11707	39.15	-120.06	1.70000
11862	40.27	-121.25	0.67500
12447	33.97	-114.49	0.87500

```
[ ]: # drop outliers, averooms 132 and 141 and AveBedrms 34 and 25
full_df = full_df.loc[full_df['AveRooms'] < 65]
```

```
[ ]: full_df[full_df['AveOccup'] > 50]
```

```
[ ]:
      MedInc  HouseAge  AveRooms  AveBedrms  Population  AveOccup  \
3364  5.5179      36.0  5.142857  1.142857     4198.0  599.714286
9172  4.2391       5.0  5.123810  0.933333     8733.0   83.171429
12104 1.6250       8.0  7.600000  0.950000     1275.0   63.750000
13034 6.1359      52.0  8.275862  1.517241     6675.0  230.172414
16420 5.7485      26.0  5.366667  0.900000     1542.0   51.400000
16669 4.2639      46.0  9.076923  1.307692     6532.0  502.461538
19006 10.2264     45.0  3.166667  0.833333     7460.0 1243.333333
```

	Latitude	Longitude	MedHouseVal
3364	40.41	-120.51	0.675
9172	34.47	-118.59	1.546
12104	33.97	-117.33	1.625

13034	38.69	-121.15	2.250
16420	37.89	-121.29	1.625
16669	35.32	-120.70	3.500
19006	38.32	-121.98	1.375

```
[ ]: # drop outliers, aveoccup 100+
full_df = full_df.loc[full_df['AveOccup'] < 100]
```

```
[ ]: full_df[full_df['Population'] > 20000]
```

	MedInc	HouseAge	AveRooms	AveBedrms	Population	AveOccup	Latitude	\
9880	2.3087	11.0	5.364518	1.059684	28566.0	4.696810	36.64	
15360	2.5729	14.0	5.270497	1.010484	35682.0	7.482072	33.35	

	Longitude	MedHouseVal
9880	-121.79	1.188
15360	-117.42	1.344

```
[ ]: # drop outliers, population 20k+
full_df = full_df.loc[full_df['Population'] < 20000]
```

1.2.5 New boxplots with significant outliers removed

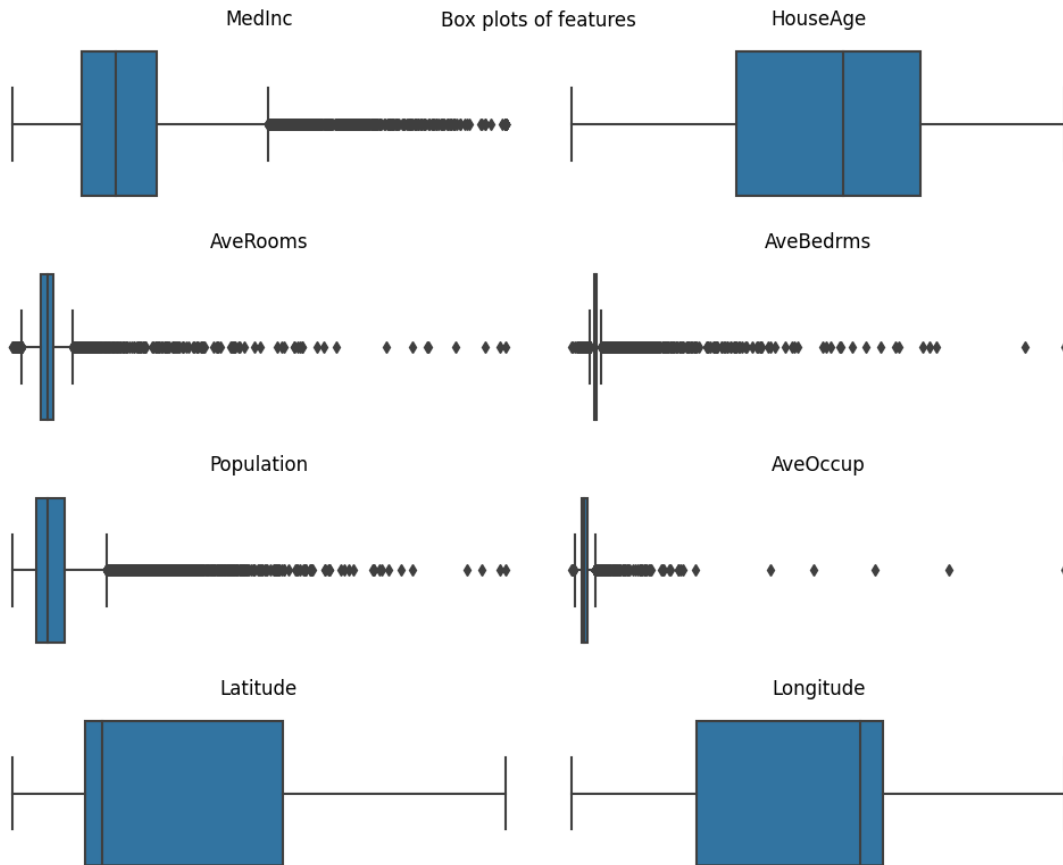
```
[ ]: fig, axes = plt.subplots(num_rows, num_cols, figsize=(10,8))
fig.subplots_adjust(hspace=0.5)

for i in range(num_features):
    row_idx = i // num_cols
    col_idx = i % num_cols

    sns.boxplot(x=full_df.iloc[:,i], ax=axes[row_idx, col_idx])
    axes[row_idx, col_idx].set_title(f'{feature_df.columns[i]}')

    if i >= num_features - (num_rows * num_cols):
        axes[row_idx, col_idx].axis('off')

plt.tight_layout()
plt.suptitle('Box plots of features')
plt.show()
```

1.3 TTS

Split the data into training and testing sets.

```
[ ]: X = full_df.drop(['MedHouseVal'], axis=1)
     y = full_df['MedHouseVal']
```

```
[ ]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
     ↪ random_state=97)
```

1.3.1 Feature and target scaling

```
[ ]: scaler_x = MinMaxScaler()

     X_train_scaled = scaler_x.fit_transform(X_train)
     X_test_scaled = scaler_x.transform(X_test)
```

```
[ ]: scaler_y = MinMaxScaler()
```

```

y_train_arr = np.array(y_train)
y_test_arr = np.array(y_test)
y_train_arr = y_train_arr.reshape(-1, 1)
y_test_arr = y_test_arr.reshape(-1, 1)

y_train_scaled = scaler_y.fit_transform(y_train_arr)
y_test_scaled = scaler_y.transform(y_test_arr)

```

1.4 Model definition

Define a deep neural network architecture for regression using Keras.

```

[ ]: def create_regression_model(optimizer='adam', activation='relu', dropout_rate=0.
    ↪2):
    model = Sequential()
    model.add(Dense(64, activation=activation, input_shape=(X_train_scaled.
    ↪shape[1],)))
    model.add(Dense(64, activation=activation))
    model.add(Dropout(dropout_rate))
    model.add(Dense(1))

    model.compile(optimizer=optimizer, loss='mean_squared_error',
    ↪metrics=['mae'])

    return model

```

```

[ ]: baseline_regression_model = create_regression_model()

```

1.5 Model training and testing

Train the model on the training set and evaluate its performance on the testing set.

1.5.1 Model training

```

[ ]: epochs = 50
    batch_size = 32

    early_stopping = EarlyStopping(monitor='val_loss', patience=5,
    ↪restore_best_weights=True)
    model_checkpoint = ModelCheckpoint('baseline_regression.h5',
    ↪save_best_only=True)

    history = baseline_regression_model.fit(X_train_scaled, y_train_scaled,
    ↪epochs=epochs, batch_size=batch_size, validation_split=0.2, verbose=1,
    ↪callbacks=[early_stopping, model_checkpoint])

```

Epoch 1/50

413/413 [=====] - 1s 1ms/step - loss: 0.0303 - mae:


0.1283 - val_loss: 0.0211 - val_mae: 0.1062
Epoch 2/50
413/413 [=====] - 0s 1ms/step - loss: 0.0235 - mae: 0.1111 - val_loss: 0.0186 - val_mae: 0.0951
Epoch 3/50
413/413 [=====] - 0s 1ms/step - loss: 0.0204 - mae: 0.1026 - val_loss: 0.0162 - val_mae: 0.0897
Epoch 4/50
413/413 [=====] - 0s 1ms/step - loss: 0.0187 - mae: 0.0975 - val_loss: 0.0153 - val_mae: 0.0865
Epoch 5/50
413/413 [=====] - 0s 1ms/step - loss: 0.0177 - mae: 0.0943 - val_loss: 0.0144 - val_mae: 0.0853
Epoch 6/50
413/413 [=====] - 0s 1ms/step - loss: 0.0172 - mae: 0.0929 - val_loss: 0.0142 - val_mae: 0.0855
Epoch 7/50
413/413 [=====] - 0s 1ms/step - loss: 0.0164 - mae: 0.0902 - val_loss: 0.0179 - val_mae: 0.1053
Epoch 8/50
413/413 [=====] - 0s 1ms/step - loss: 0.0163 - mae: 0.0903 - val_loss: 0.0137 - val_mae: 0.0815
Epoch 9/50
413/413 [=====] - 1s 1ms/step - loss: 0.0157 - mae: 0.0887 - val_loss: 0.0131 - val_mae: 0.0820
Epoch 10/50
413/413 [=====] - 0s 1ms/step - loss: 0.0155 - mae: 0.0875 - val_loss: 0.0140 - val_mae: 0.0884
Epoch 11/50
413/413 [=====] - 0s 1ms/step - loss: 0.0153 - mae: 0.0871 - val_loss: 0.0130 - val_mae: 0.0805
Epoch 12/50
413/413 [=====] - 0s 1ms/step - loss: 0.0154 - mae: 0.0875 - val_loss: 0.0146 - val_mae: 0.0826
Epoch 13/50
413/413 [=====] - 0s 1ms/step - loss: 0.0151 - mae: 0.0864 - val_loss: 0.0131 - val_mae: 0.0795
Epoch 14/50
413/413 [=====] - 0s 1ms/step - loss: 0.0150 - mae: 0.0860 - val_loss: 0.0128 - val_mae: 0.0788
Epoch 15/50
413/413 [=====] - 0s 1ms/step - loss: 0.0147 - mae: 0.0854 - val_loss: 0.0144 - val_mae: 0.0817
Epoch 16/50
413/413 [=====] - 0s 1ms/step - loss: 0.0146 - mae: 0.0848 - val_loss: 0.0130 - val_mae: 0.0829
Epoch 17/50
413/413 [=====] - 1s 1ms/step - loss: 0.0148 - mae:

0.0854 - val_loss: 0.0133 - val_mae: 0.0776
 Epoch 18/50
 413/413 [=====] - 1s 1ms/step - loss: 0.0145 - mae: 0.0842 - val_loss: 0.0156 - val_mae: 0.0925
 Epoch 19/50
 413/413 [=====] - 1s 1ms/step - loss: 0.0143 - mae: 0.0837 - val_loss: 0.0123 - val_mae: 0.0772
 Epoch 20/50
 413/413 [=====] - 0s 1ms/step - loss: 0.0145 - mae: 0.0844 - val_loss: 0.0127 - val_mae: 0.0811
 Epoch 21/50
 413/413 [=====] - 0s 1ms/step - loss: 0.0142 - mae: 0.0841 - val_loss: 0.0122 - val_mae: 0.0770
 Epoch 22/50
 413/413 [=====] - 1s 1ms/step - loss: 0.0142 - mae: 0.0835 - val_loss: 0.0130 - val_mae: 0.0772
 Epoch 23/50
 413/413 [=====] - 0s 1ms/step - loss: 0.0142 - mae: 0.0836 - val_loss: 0.0123 - val_mae: 0.0763
 Epoch 24/50
 413/413 [=====] - 0s 1ms/step - loss: 0.0141 - mae: 0.0830 - val_loss: 0.0126 - val_mae: 0.0766
 Epoch 25/50
 413/413 [=====] - 0s 1ms/step - loss: 0.0142 - mae: 0.0837 - val_loss: 0.0123 - val_mae: 0.0779
 Epoch 26/50
 413/413 [=====] - 0s 1ms/step - loss: 0.0138 - mae: 0.0824 - val_loss: 0.0122 - val_mae: 0.0757
 Epoch 27/50
 413/413 [=====] - 0s 1ms/step - loss: 0.0138 - mae: 0.0822 - val_loss: 0.0122 - val_mae: 0.0758
 Epoch 28/50
 413/413 [=====] - 0s 1ms/step - loss: 0.0139 - mae: 0.0824 - val_loss: 0.0136 - val_mae: 0.0788
 Epoch 29/50
 413/413 [=====] - 0s 1ms/step - loss: 0.0138 - mae: 0.0826 - val_loss: 0.0122 - val_mae: 0.0761
 Epoch 30/50
 413/413 [=====] - 0s 1ms/step - loss: 0.0140 - mae: 0.0826 - val_loss: 0.0120 - val_mae: 0.0746
 Epoch 31/50
 413/413 [=====] - 0s 1ms/step - loss: 0.0135 - mae: 0.0814 - val_loss: 0.0134 - val_mae: 0.0853
 Epoch 32/50
 413/413 [=====] - 0s 1ms/step - loss: 0.0136 - mae: 0.0820 - val_loss: 0.0150 - val_mae: 0.0823
 Epoch 33/50
 413/413 [=====] - 0s 1ms/step - loss: 0.0135 - mae:


0.0815 - val_loss: 0.0117 - val_mae: 0.0743
 Epoch 34/50
 413/413 [=====] - 0s 1ms/step - loss: 0.0136 - mae: 0.0818 - val_loss: 0.0117 - val_mae: 0.0762
 Epoch 35/50
 413/413 [=====] - 1s 1ms/step - loss: 0.0135 - mae: 0.0814 - val_loss: 0.0123 - val_mae: 0.0804
 Epoch 36/50
 413/413 [=====] - 0s 1ms/step - loss: 0.0134 - mae: 0.0808 - val_loss: 0.0116 - val_mae: 0.0745
 Epoch 37/50
 413/413 [=====] - 0s 1ms/step - loss: 0.0135 - mae: 0.0810 - val_loss: 0.0115 - val_mae: 0.0744
 Epoch 38/50
 413/413 [=====] - 0s 1ms/step - loss: 0.0134 - mae: 0.0808 - val_loss: 0.0119 - val_mae: 0.0762
 Epoch 39/50
 413/413 [=====] - 0s 1ms/step - loss: 0.0132 - mae: 0.0805 - val_loss: 0.0115 - val_mae: 0.0753
 Epoch 40/50
 413/413 [=====] - 0s 1ms/step - loss: 0.0134 - mae: 0.0808 - val_loss: 0.0125 - val_mae: 0.0768
 Epoch 41/50
 413/413 [=====] - 0s 1ms/step - loss: 0.0133 - mae: 0.0810 - val_loss: 0.0122 - val_mae: 0.0801
 Epoch 42/50
 413/413 [=====] - 0s 1ms/step - loss: 0.0131 - mae: 0.0801 - val_loss: 0.0115 - val_mae: 0.0741
 Epoch 43/50
 413/413 [=====] - 0s 1ms/step - loss: 0.0133 - mae: 0.0802 - val_loss: 0.0131 - val_mae: 0.0766
 Epoch 44/50
 413/413 [=====] - 0s 1ms/step - loss: 0.0132 - mae: 0.0811 - val_loss: 0.0131 - val_mae: 0.0760
 Epoch 45/50
 413/413 [=====] - 0s 1ms/step - loss: 0.0131 - mae: 0.0803 - val_loss: 0.0120 - val_mae: 0.0780
 Epoch 46/50
 413/413 [=====] - 0s 1ms/step - loss: 0.0129 - mae: 0.0794 - val_loss: 0.0116 - val_mae: 0.0746
 Epoch 47/50
 413/413 [=====] - 0s 1ms/step - loss: 0.0130 - mae: 0.0796 - val_loss: 0.0115 - val_mae: 0.0758

1.5.2 Baseline training perf

```
[ ]: def training_plots(history, metric):  
    """  
    Function to plot training and validation history for a given metric  
    inputs: history - model training history  
           metric - desired metric for plotting  
    returns: -  
    """  
    epochs = len(history.history[f'{metric}'])  
    plt.figure(figsize=(8, 6))  
    sns.lineplot(x=range(1, epochs + 1), y=history.history[f'{metric}'],  
↳label=f'Training {metric}')  
    sns.lineplot(x=range(1, epochs + 1), y=history.history[f'val_{metric}'],  
↳label=f'Validation {metric}')  
    plt.xlabel('Epoch')  
    plt.ylabel(f'{metric}')  
    plt.title(f'Training and Validation {metric}')  
    plt.legend()  
    plt.show()  
  
[ ]: training_plots(history, 'mae')
```



```
[ ]: training_plots(history, 'loss')
```



1.5.3 Basemodel performance in testing

```
[ ]: loss, mae = baseline_regression_model.evaluate(X_test_scaled, y_test_scaled, verbose=0)
print(f'Test Loss: {loss:.4f}')
print(f'Test MAE: {mae:.4f}')
```

Test Loss: 0.0123


Test MAE: 0.0757

```
[ ]: y_pred = baseline_regression_model.predict(X_test_scaled)
y_pred_inv = scaler_y.inverse_transform(y_pred)
```

129/129 [=====] - 0s 571us/step

```
[ ]: def testing_plot(y_true, y_pred, tuned=0):
    """
    Function to scatterplot testing performance of a regression model
    Inputs: y_true - Actual targets
            y_pred - Model predicted targets
            tuned - parameter to allow for plotting of 2 scatterplots
    Returns: -
    """
    if tuned:
        plt.scatter(y_true, y_pred, color='green', label='Tuned model')
    else:
        plt.scatter(y_true, y_pred, color='blue', label='Baseline model')
    plt.xlabel('True values')
    plt.ylabel('Predicted values')
    plt.title('True vs Pred')
    plt.plot([min(y_true), max(y_true)], [min(y_true), max(y_true)],
             color='red')

[ ]: testing_plot(y_test, y_pred_inv)
plt.show()
```



1.6 Parameter tuning

Tune the hyperparameters of the model to achieve better performance (e.g., number of hidden layers, activation functions, learning rate, number of epochs, etc.).

```
[ ]: tuning_regression_model = KerasRegressor(build_fn=create_regression_model,
↳ verbose=0)

param_grid = {
    'optimizer': ['adam', 'sgd'],
    'activation': ['relu', 'tanh'],
    'batch_size': [16, 32, 64],
    'epochs': [20, 50],
    'dropout_rate': [0.2, 0.3]
}
```

```
C:\Users\Reed Oken\AppData\Local\Temp\ipykernel_16472\1470353211.py:1:
DeprecationWarning: KerasRegressor is deprecated, use Sci-Keras
(https://github.com/adriangb/scikeras) instead. See
https://www.adriangb.com/scikeras/stable/migration.html for help migrating.
    tuning_regression_model = KerasRegressor(build_fn=create_regression_model,
    verbose=0)
```

```
[ ]: regressor_grid = GridSearchCV(estimator=tuning_regression_model,
↳ param_grid=param_grid, cv=3, n_jobs=-1)

grid_result = regressor_grid.fit(X_train_scaled, y_train_scaled,
↳ callbacks=[early_stopping])
```

```
WARNING:tensorflow:Early stopping conditioned on metric `val_loss` which is not
available. Available metrics are: loss,mae
WARNING:tensorflow:Early stopping conditioned on metric `val_loss` which is not
available. Available metrics are: loss,mae
WARNING:tensorflow:Early stopping conditioned on metric `val_loss` which is not
available. Available metrics are: loss,mae
WARNING:tensorflow:Early stopping conditioned on metric `val_loss` which is not
available. Available metrics are: loss,mae
WARNING:tensorflow:Early stopping conditioned on metric `val_loss` which is not
available. Available metrics are: loss,mae
WARNING:tensorflow:Early stopping conditioned on metric `val_loss` which is not
available. Available metrics are: loss,mae
WARNING:tensorflow:Early stopping conditioned on metric `val_loss` which is not
available. Available metrics are: loss,mae
WARNING:tensorflow:Early stopping conditioned on metric `val_loss` which is not
available. Available metrics are: loss,mae
WARNING:tensorflow:Early stopping conditioned on metric `val_loss` which is not
available. Available metrics are: loss,mae
WARNING:tensorflow:Early stopping conditioned on metric `val_loss` which is not
available. Available metrics are: loss,mae
```

[illegible]

```

WARNING:tensorflow:Early stopping conditioned on metric `val_loss` which is not
available. Available metrics are: loss,mae
WARNING:tensorflow:Early stopping conditioned on metric `val_loss` which is not
available. Available metrics are: loss,mae
WARNING:tensorflow:Early stopping conditioned on metric `val_loss` which is not
available. Available metrics are: loss,mae
WARNING:tensorflow:Early stopping conditioned on metric `val_loss` which is not
available. Available metrics are: loss,mae
WARNING:tensorflow:Early stopping conditioned on metric `val_loss` which is not
available. Available metrics are: loss,mae
WARNING:tensorflow:Early stopping conditioned on metric `val_loss` which is not
available. Available metrics are: loss,mae
WARNING:tensorflow:Early stopping conditioned on metric `val_loss` which is not
available. Available metrics are: loss,mae
WARNING:tensorflow:Early stopping conditioned on metric `val_loss` which is not
available. Available metrics are: loss,mae
WARNING:tensorflow:Early stopping conditioned on metric `val_loss` which is not
available. Available metrics are: loss,mae
WARNING:tensorflow:Early stopping conditioned on metric `val_loss` which is not
available. Available metrics are: loss,mae
WARNING:tensorflow:Early stopping conditioned on metric `val_loss` which is not
available. Available metrics are: loss,mae
WARNING:tensorflow:Early stopping conditioned on metric `val_loss` which is not
available. Available metrics are: loss,mae
WARNING:tensorflow:Early stopping conditioned on metric `val_loss` which is not
available. Available metrics are: loss,mae
WARNING:tensorflow:Early stopping conditioned on metric `val_loss` which is not
available. Available metrics are: loss,mae
WARNING:tensorflow:Early stopping conditioned on metric `val_loss` which is not
available. Available metrics are: loss,mae
WARNING:tensorflow:Early stopping conditioned on metric `val_loss` which is not
available. Available metrics are: loss,mae

```

```

[ ]: gs_regression_df = pd.DataFrame(grid_result.cv_results_)
gs_regression_df = gs_regression_df.drop(columns=['mean_fit_time',
↳ 'std_fit_time', 'mean_score_time', 'std_score_time', 'params',
↳ 'split0_test_score', 'split1_test_score', 'split2_test_score'])

gs_regression_df = gs_regression_df.sort_values(by='rank_test_score',
↳ ascending=True)
gs_regression_df.head(20).round(4)

```

```

[ ]:   param_activation param_batch_size param_dropout_rate param_epochs \
10          relu                32                0.2             50
22          relu                64                0.3             50
6          relu                16                0.3             50
14         relu                32                0.3             50

```

18	relu	64	0.2	50
2	relu	16	0.2	50
8	relu	32	0.2	20
0	relu	16	0.2	20
12	relu	32	0.3	20
16	relu	64	0.2	20
4	relu	16	0.3	20
20	relu	64	0.3	20
30	tanh	16	0.3	50
26	tanh	16	0.2	50
38	tanh	32	0.3	50
34	tanh	32	0.2	50
3	relu	16	0.2	50
7	relu	16	0.3	50
28	tanh	16	0.3	20
42	tanh	64	0.2	50

	param_optimizer	mean_test_score	std_test_score	rank_test_score
10	adam	-0.0125	0.0006	1
22	adam	-0.0126	0.0004	2
6	adam	-0.0128	0.0003	3
14	adam	-0.0129	0.0007	4
18	adam	-0.0130	0.0005	5
2	adam	-0.0134	0.0006	6
8	adam	-0.0136	0.0008	7
0	adam	-0.0137	0.0016	8
12	adam	-0.0140	0.0011	9
16	adam	-0.0142	0.0008	10
4	adam	-0.0143	0.0011	11
20	adam	-0.0144	0.0009	12
30	adam	-0.0165	0.0010	13
26	adam	-0.0168	0.0006	14
38	adam	-0.0170	0.0008	15
34	adam	-0.0177	0.0009	16
3	sgd	-0.0192	0.0013	17
7	sgd	-0.0194	0.0007	18
28	adam	-0.0196	0.0008	19
42	adam	-0.0196	0.0003	20

1.7 Model evaluation

Compare the performance of the tuned model with the baseline model (i.e., the initial model without any hyperparameter tuning).

```
[ ]: # best params identified via grid search
activation = 'relu'
batch_size = 16
```

```
dropout_rate = 0.2
epochs = 50
optimizer = 'adam'
```

```
[ ]: tuned_regression_model = create_regression_model(activation=activation,
↳optimizer=optimizer, dropout_rate=dropout_rate)
```

```
[ ]: model_checkpoint = ModelCheckpoint('tuned_regression.h5', save_best_only=True)

history = tuned_regression_model.fit(X_train_scaled, y_train_scaled,
↳epochs=epochs, batch_size=batch_size, validation_split=0.2, verbose=1,
↳callbacks=[early_stopping, model_checkpoint])
```

Epoch 1/50

826/826 [=====] - 1s 1ms/step - loss: 0.0295 - mae: 0.1259 - val_loss: 0.0203 - val_mae: 0.1015

Epoch 2/50

826/826 [=====] - 1s 1ms/step - loss: 0.0221 - mae: 0.1071 - val_loss: 0.0238 - val_mae: 0.1037

Epoch 3/50

826/826 [=====] - 1s 1ms/step - loss: 0.0192 - mae: 0.0988 - val_loss: 0.0159 - val_mae: 0.0864

Epoch 4/50

826/826 [=====] - 1s 1ms/step - loss: 0.0177 - mae: 0.0947 - val_loss: 0.0147 - val_mae: 0.0903

Epoch 5/50

826/826 [=====] - 1s 1ms/step - loss: 0.0172 - mae: 0.0933 - val_loss: 0.0145 - val_mae: 0.0871

Epoch 6/50

826/826 [=====] - 1s 1ms/step - loss: 0.0165 - mae: 0.0908 - val_loss: 0.0138 - val_mae: 0.0850

Epoch 7/50

826/826 [=====] - 1s 1ms/step - loss: 0.0161 - mae: 0.0897 - val_loss: 0.0145 - val_mae: 0.0834

Epoch 8/50

826/826 [=====] - 1s 1ms/step - loss: 0.0161 - mae: 0.0896 - val_loss: 0.0135 - val_mae: 0.0847

Epoch 9/50

826/826 [=====] - 1s 1ms/step - loss: 0.0155 - mae: 0.0879 - val_loss: 0.0129 - val_mae: 0.0822

Epoch 10/50

826/826 [=====] - 1s 1ms/step - loss: 0.0153 - mae: 0.0874 - val_loss: 0.0128 - val_mae: 0.0786

Epoch 11/50

826/826 [=====] - 1s 1ms/step - loss: 0.0151 - mae: 0.0868 - val_loss: 0.0141 - val_mae: 0.0815

Epoch 12/50

826/826 [=====] - 1s 1ms/step - loss: 0.0152 - mae:


```

0.0866 - val_loss: 0.0150 - val_mae: 0.0833
Epoch 13/50
826/826 [=====] - 1s 1ms/step - loss: 0.0153 - mae:
0.0871 - val_loss: 0.0135 - val_mae: 0.0787
Epoch 14/50
826/826 [=====] - 1s 1ms/step - loss: 0.0148 - mae:
0.0855 - val_loss: 0.0125 - val_mae: 0.0787
Epoch 15/50
826/826 [=====] - 1s 998us/step - loss: 0.0145 - mae:
0.0847 - val_loss: 0.0123 - val_mae: 0.0773
Epoch 16/50
826/826 [=====] - 1s 998us/step - loss: 0.0146 - mae:
0.0850 - val_loss: 0.0124 - val_mae: 0.0787
Epoch 17/50
826/826 [=====] - 1s 990us/step - loss: 0.0145 - mae:
0.0843 - val_loss: 0.0125 - val_mae: 0.0770
Epoch 18/50
826/826 [=====] - 1s 1ms/step - loss: 0.0146 - mae:
0.0848 - val_loss: 0.0124 - val_mae: 0.0783
Epoch 19/50
826/826 [=====] - 1s 1ms/step - loss: 0.0144 - mae:
0.0841 - val_loss: 0.0130 - val_mae: 0.0833
Epoch 20/50
826/826 [=====] - 1s 1ms/step - loss: 0.0145 - mae:
0.0847 - val_loss: 0.0123 - val_mae: 0.0800


```

1.7.1 Tuned model training perf

```
[ ]: training_plots(history, 'mae')
```



```
[ ]: training_plots(history, 'loss')
```



1.7.2 Tuned model testing perf

```
[ ]: baseline_loss, baseline_mae = baseline_regression_model.evaluate(X_test_scaled,
    ↪ y_test_scaled, verbose=0)
    tuned_loss, tuned_mae = tuned_regression_model.evaluate(X_test_scaled,
    ↪ y_test_scaled, verbose=0)

print(f'Baseline model test loss: {baseline_loss:.4f}')
print(f'Tuned model test loss: {tuned_loss:.4f}\n')

print(f'Baseline model test MAE: {baseline_mae:.4f}')
print(f'Tuned model test MAE: {tuned_mae:.4f}')
```

Baseline model test loss: 0.0123

Tuned model test loss: 0.0129

Baseline model test MAE: 0.0757

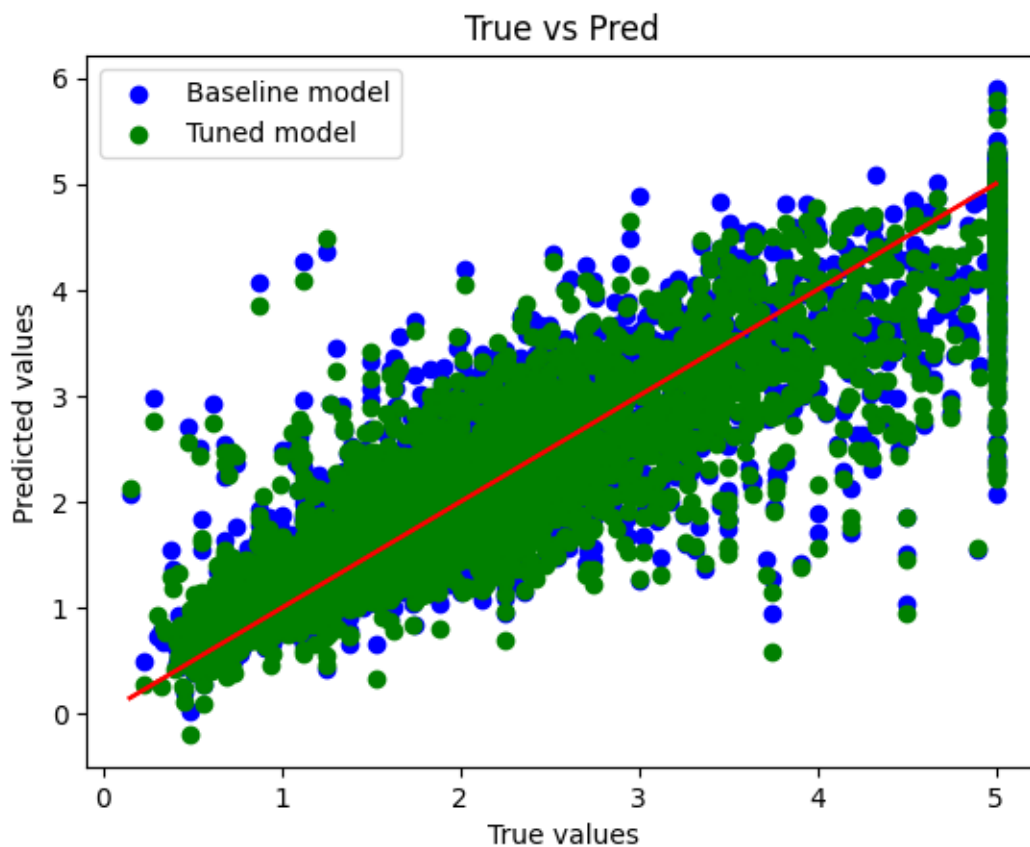
Tuned model test MAE: 0.0787


```
[ ]: y_pred_baseline = baseline_regression_model.predict(X_test_scaled)
      y_pred_tuned = tuned_regression_model.predict(X_test_scaled)

      y_pred_baseline_inv = scaler_y.inverse_transform(y_pred_baseline)
      y_pred_tuned_inv = scaler_y.inverse_transform(y_pred_tuned)
```

```
1/129 [...] - ETA: 1s129/129
[=====] - 0s 643us/step
129/129 [=====] - 0s 585us/step
```

```
[ ]: testing_plot(y_test, y_pred_baseline_inv)
      testing_plot(y_test, y_pred_tuned_inv, tuned=1)
      plt.legend()
      plt.show()
```



2 Classification problem

2.1 Load data

Load the Boston Housing dataset from the Keras library.

No new data will be loaded as the data loaded for the regression problem is the same as the data required for the classification problem.

2.2 EDA

Explore and preprocess the data (e.g., normalization, one-hot encoding, etc.).

EDA and preprocessing completed within the scope of the regression setup.

2.3 Binary target


Convert the target variable into a binary variable (i.e., expensive or not expensive).

2.3.1 Investigate the ‘breakpoint’ for expensive vs non-expensive

```
[ ]: full_df['MedHouseVal'].describe()
```

```
[ ]: count      20632.000000
      mean        2.068538
      std         1.153873
      min         0.149990
      25%         1.196000
      50%         1.797000
      75%         2.647250
      max         5.000010
      Name: MedHouseVal, dtype: float64
```

```
[ ]: sns.boxplot(x=full_df['MedHouseVal'])
      plt.show()
```



Expensive houses will be $\text{MedHouseVal} > 3$

```
[ ]: # create new binary target 'expensive'
expensive = 3

full_df['expensive'] = full_df['MedHouseVal'].apply(lambda x: 1 if x > 3
↪expensive else 0)
```

2.4 TTS

Split the data into training and testing sets.

```
[ ]: X = full_df.drop(['MedHouseVal', 'expensive'], axis=1)
y = full_df['expensive']

[ ]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↪random_state=97)
```

2.4.1 Feature scaling

```
[ ]: scaler_x = MinMaxScaler()

X_train_scaled = scaler_x.fit_transform(X_train)
X_test_scaled = scaler_x.transform(X_test)
```

2.5 Define model

Define a deep neural network architecture for classification using Keras.

```
[ ]: def create_classification_model(optimizer='adam', activation='relu',
    ↪ dropout_rate=0.2):
    """
    Function to create a DNN classifier
    Inputs: Optimizer - model optimizer, default: adam
            actiavation - activation function for hidden layers, default: relu
            dropout_rate - dropout rate for hidden layer, default: 0.2
    Returns: compiled model
    """
    model = Sequential()
    model.add(Dense(64, activation=activation, input_shape=(X_train_scaled.
    ↪ shape[1],)))
    model.add(Dense(64, activation=activation))
    model.add(Dropout(dropout_rate))
    model.add(Dense(1, activation='sigmoid'))

    model.compile(optimizer=optimizer, loss='binary_crossentropy',
    ↪ metrics=['accuracy'])

    return model
```

2.6 Train and test model

Train the model on the training set and evaluate its performance on the testing set.

```
[ ]: baseline_classification_model = create_classification_model()

[ ]: epochs = 50
    batch_size = 32

    early_stopping = EarlyStopping(monitor='val_loss', patience=5,
    ↪ restore_best_weights=True)
    model_checkpoint = ModelCheckpoint('baseline_regression.h5',
    ↪ save_best_only=True)
```

```
history = baseline_classification_model.fit(X_train_scaled, y_train,  
↳ epochs=epochs, batch_size=batch_size, validation_split=0.2, verbose=1,  
↳ callbacks=[early_stopping, model_checkpoint])
```

Epoch 1/50

```
c:\Users\Reed Oken\AppData\Local\Programs\Python\Python310\lib\site-  
packages\keras\engine\data_adapter.py:1700: FutureWarning: The behavior of  
`series[i:j]` with an integer-dtype index is deprecated. In a future version,  
this will be treated as *label-based* indexing, consistent with e.g. `series[i]`  
lookups. To retain the old behavior, use `series.iloc[i:j]`. To get the future  
behavior, use `series.loc[i:j]`.
```

```
    return t[start:end]
```

```
413/413 [=====] - 1s 1ms/step - loss: 0.3683 -  
accuracy: 0.8509 - val_loss: 0.2877 - val_accuracy: 0.8843
```

Epoch 2/50

```
413/413 [=====] - 0s 1ms/step - loss: 0.2974 -  
accuracy: 0.8807 - val_loss: 0.2733 - val_accuracy: 0.8949
```

Epoch 3/50

```
413/413 [=====] - 0s 1ms/step - loss: 0.2912 -  
accuracy: 0.8828 - val_loss: 0.2638 - val_accuracy: 0.8964
```

Epoch 4/50

```
413/413 [=====] - 1s 1ms/step - loss: 0.2817 -  
accuracy: 0.8856 - val_loss: 0.2614 - val_accuracy: 0.8973
```

Epoch 5/50

```
413/413 [=====] - 0s 1ms/step - loss: 0.2763 -  
accuracy: 0.8863 - val_loss: 0.2476 - val_accuracy: 0.9003
```

Epoch 6/50

```
413/413 [=====] - 1s 1ms/step - loss: 0.2680 -  
accuracy: 0.8912 - val_loss: 0.2440 - val_accuracy: 0.9018
```

Epoch 7/50

```
413/413 [=====] - 0s 1ms/step - loss: 0.2595 -  
accuracy: 0.8919 - val_loss: 0.2412 - val_accuracy: 0.9006
```

Epoch 8/50

```
413/413 [=====] - 0s 1ms/step - loss: 0.2562 -  
accuracy: 0.8944 - val_loss: 0.2488 - val_accuracy: 0.9000
```

Epoch 9/50

```
413/413 [=====] - 0s 1ms/step - loss: 0.2553 -  
accuracy: 0.8950 - val_loss: 0.2316 - val_accuracy: 0.9085
```

Epoch 10/50

```
413/413 [=====] - 0s 1ms/step - loss: 0.2510 -  
accuracy: 0.8979 - val_loss: 0.2493 - val_accuracy: 0.9028
```

Epoch 11/50

```
413/413 [=====] - 0s 1ms/step - loss: 0.2509 -  
accuracy: 0.8964 - val_loss: 0.2358 - val_accuracy: 0.9061
```

Epoch 12/50

```
413/413 [=====] - 1s 1ms/step - loss: 0.2492 -  
accuracy: 0.8972 - val_loss: 0.2285 - val_accuracy: 0.9100
```

Epoch 13/50
413/413 [=====] - 0s 1ms/step - loss: 0.2458 -
accuracy: 0.8986 - val_loss: 0.2247 - val_accuracy: 0.9097
Epoch 14/50
413/413 [=====] - 0s 1ms/step - loss: 0.2438 -
accuracy: 0.8979 - val_loss: 0.2248 - val_accuracy: 0.9097
Epoch 15/50
413/413 [=====] - 0s 1ms/step - loss: 0.2448 -
accuracy: 0.8997 - val_loss: 0.2377 - val_accuracy: 0.9067
Epoch 16/50
413/413 [=====] - 0s 1ms/step - loss: 0.2449 -
accuracy: 0.8977 - val_loss: 0.2227 - val_accuracy: 0.9146
Epoch 17/50
413/413 [=====] - 0s 1ms/step - loss: 0.2413 -
accuracy: 0.9026 - val_loss: 0.2246 - val_accuracy: 0.9121
Epoch 18/50
413/413 [=====] - 0s 1ms/step - loss: 0.2427 -
accuracy: 0.9006 - val_loss: 0.2259 - val_accuracy: 0.9115
Epoch 19/50
413/413 [=====] - 0s 1ms/step - loss: 0.2398 -
accuracy: 0.9011 - val_loss: 0.2218 - val_accuracy: 0.9121
Epoch 20/50
413/413 [=====] - 1s 1ms/step - loss: 0.2385 -
accuracy: 0.9028 - val_loss: 0.2201 - val_accuracy: 0.9131
Epoch 21/50
413/413 [=====] - 0s 1ms/step - loss: 0.2380 -
accuracy: 0.9037 - val_loss: 0.2194 - val_accuracy: 0.9155
Epoch 22/50
413/413 [=====] - 0s 1ms/step - loss: 0.2374 -
accuracy: 0.9038 - val_loss: 0.2213 - val_accuracy: 0.9143
Epoch 23/50
413/413 [=====] - 0s 1ms/step - loss: 0.2399 -
accuracy: 0.9004 - val_loss: 0.2193 - val_accuracy: 0.9152
Epoch 24/50
413/413 [=====] - 0s 1ms/step - loss: 0.2387 -
accuracy: 0.9009 - val_loss: 0.2231 - val_accuracy: 0.9103
Epoch 25/50
413/413 [=====] - 0s 1ms/step - loss: 0.2386 -
accuracy: 0.9020 - val_loss: 0.2202 - val_accuracy: 0.9131
Epoch 26/50
413/413 [=====] - 0s 1ms/step - loss: 0.2352 -
accuracy: 0.9017 - val_loss: 0.2208 - val_accuracy: 0.9134
Epoch 27/50
413/413 [=====] - 0s 1ms/step - loss: 0.2365 -
accuracy: 0.9016 - val_loss: 0.2179 - val_accuracy: 0.9146
Epoch 28/50
413/413 [=====] - 0s 1ms/step - loss: 0.2333 -
accuracy: 0.9050 - val_loss: 0.2193 - val_accuracy: 0.9137


```

Epoch 29/50
413/413 [=====] - 0s 1ms/step - loss: 0.2359 -
accuracy: 0.9047 - val_loss: 0.2209 - val_accuracy: 0.9158
Epoch 30/50
413/413 [=====] - 0s 1ms/step - loss: 0.2337 -
accuracy: 0.9040 - val_loss: 0.2225 - val_accuracy: 0.9125
Epoch 31/50
413/413 [=====] - 0s 1ms/step - loss: 0.2351 -
accuracy: 0.9031 - val_loss: 0.2164 - val_accuracy: 0.9176
Epoch 32/50
413/413 [=====] - 0s 1ms/step - loss: 0.2293 -
accuracy: 0.9047 - val_loss: 0.2148 - val_accuracy: 0.9161
Epoch 33/50
413/413 [=====] - 0s 1ms/step - loss: 0.2310 -
accuracy: 0.9040 - val_loss: 0.2164 - val_accuracy: 0.9143
Epoch 34/50
413/413 [=====] - 0s 1ms/step - loss: 0.2303 -
accuracy: 0.9043 - val_loss: 0.2160 - val_accuracy: 0.9137
Epoch 35/50
413/413 [=====] - 0s 1ms/step - loss: 0.2303 -
accuracy: 0.9055 - val_loss: 0.2301 - val_accuracy: 0.9091
Epoch 36/50
413/413 [=====] - 0s 1ms/step - loss: 0.2275 -
accuracy: 0.9051 - val_loss: 0.2166 - val_accuracy: 0.9146
Epoch 37/50
413/413 [=====] - 0s 1ms/step - loss: 0.2287 -
accuracy: 0.9053 - val_loss: 0.2210 - val_accuracy: 0.9091


```

2.6.1 Baseline training progression plots

```
[ ]: training_plots(history, 'accuracy')
```



```
[ ]: training_plots(history, 'loss')
```

2.6.2 Baseline testing performance

```
[ ]: loss, accuracy = baseline_classification_model.evaluate(X_test_scaled, y_test, verbose=0)
print(f'Test Loss: {loss:.4f}')
print(f'Test accuracy: {accuracy:.4f}')
```

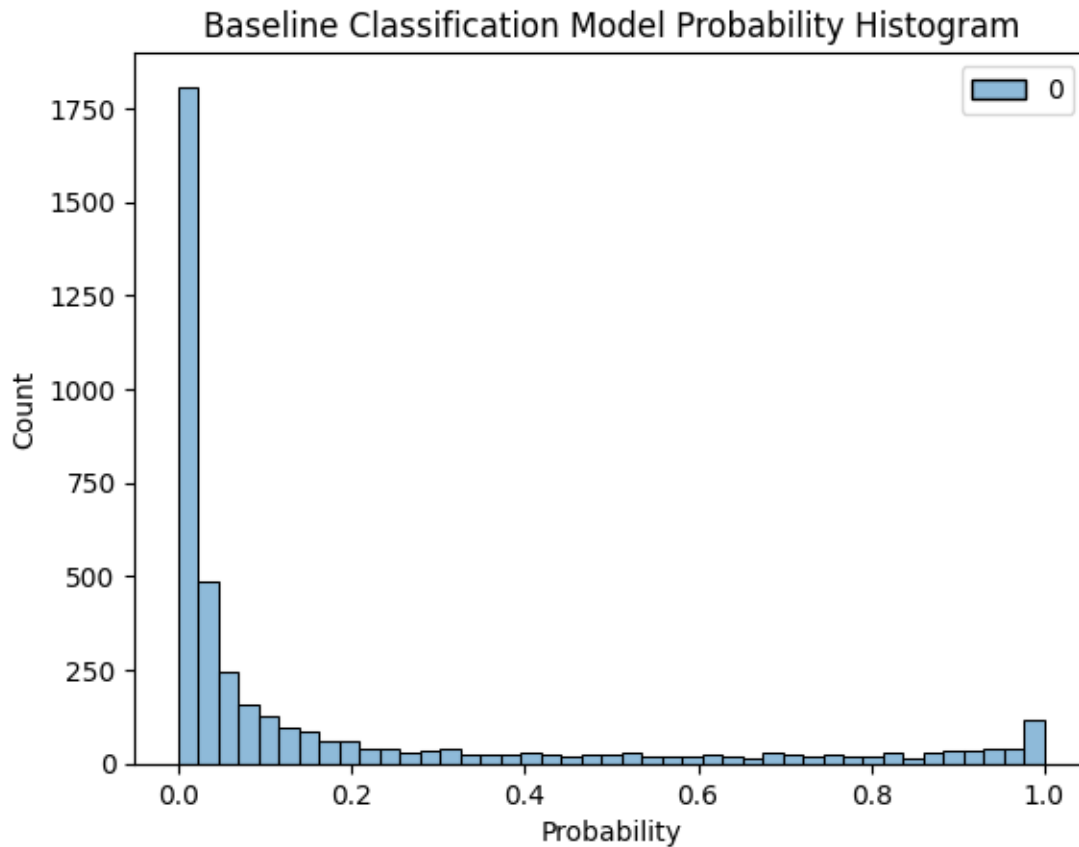
Test Loss: 0.2132

Test accuracy: 0.9159

```
[ ]: predictions = baseline_classification_model.predict(X_test_scaled)
threshold = 0.5
y_pred_baseline = np.where(predictions >= threshold, 1, 0)
```

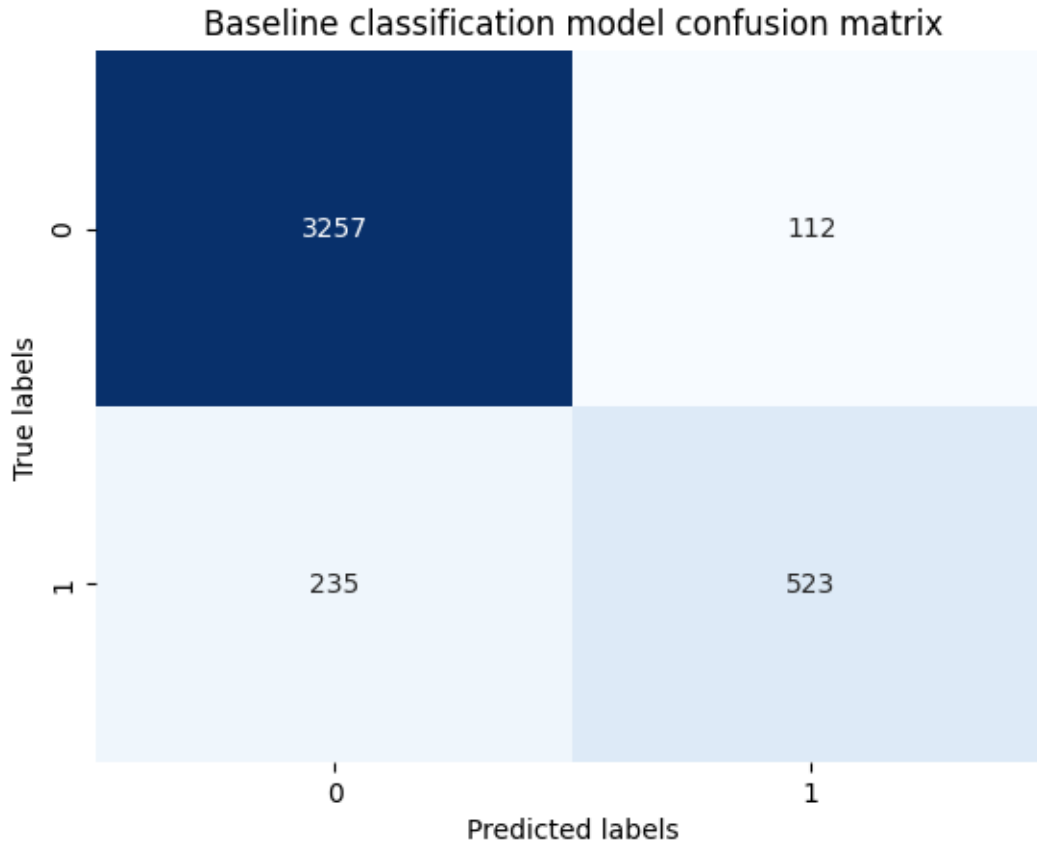
129/129 [=====] - 0s 580us/step

```
[ ]: sns.histplot(predictions)
plt.title('Baseline Classification Model Probability Histogram')
plt.xlabel('Probability')
plt.show()
```



```
[ ]: confusion_mat = confusion_matrix(y_test, y_pred_baseline)

sns.heatmap(confusion_mat, annot=True, cmap='Blues', fmt='g', cbar=False)
plt.xlabel('Predicted labels')
plt.ylabel('True labels')
plt.title('Baseline classification model confusion matrix')
plt.show()
```



```
[ ]: baseline_acc = accuracy_score(y_test, y_pred_baseline)
baseline_f1 = f1_score(y_test, y_pred_baseline)
baseline_precision = precision_score(y_test, y_pred_baseline)
baseline_recall = recall_score(y_test, y_pred_baseline)

[ ]: print(f'Baseline classification model test accuracy: {baseline_acc:.4f}')
print(f'Baseline classification model test precision: {baseline_precision:.4f}')
print(f'Baseline classification model test recall: {baseline_recall:.4f}')
print(f'Baseline classification model test F1: {baseline_f1:.4f}')
```

```
Baseline classification model test accuracy: 0.9159
Baseline classification model test precision: 0.8236
Baseline classification model test recall: 0.6900
Baseline classification model test F1: 0.7509
```

2.7 Parameter tuning

Tune the hyperparameters of the model to achieve better performance (e.g., number of hidden layers, activation functions, learning rate, number of epochs, etc.).

```
[ ]: tuning_classification_model = ↳ KerasClassifier(build_fn=create_classification_model, verbose=0)
```

C:\Users\Reed Oken\AppData\Local\Temp\ipykernel_16472\3533025257.py:1:
 DeprecationWarning: KerasClassifier is deprecated, use Sci-Keras
 (https://github.com/adriangb/scikeras) instead. See
 https://www.adriangb.com/scikeras/stable/migration.html for help migrating.
 tuning_classification_model =
 KerasClassifier(build_fn=create_classification_model, verbose=0)

```
[ ]: param_grid = {
    'optimizer': ['adam', 'sgd'],
    'activation': ['relu', 'tanh'],
    'batch_size': [16, 32, 64],
    'epochs': [20, 50],
    'dropout_rate': [0.2, 0.3]
}
```

```
[ ]: classifier_grid = GridSearchCV(estimator=tuning_classification_model, ↳
    param_grid=param_grid, cv=3, n_jobs=-1)

model_checkpoint = ModelCheckpoint('tuning_classification.h5', ↳
    save_best_only=True)
```

```
[ ]: grid_result = classifier_grid.fit(X_train_scaled, y_train, ↳
    callbacks=[early_stopping], validation_split=0.2)
```

```
[ ]: gs_classification_df = pd.DataFrame(grid_result.cv_results_)
gs_classification_df = gs_classification_df.drop(columns=['mean_fit_time', ↳
    'std_fit_time', 'mean_score_time', 'std_score_time', 'params', ↳
    'split0_test_score', 'split1_test_score', 'split2_test_score'])

gs_classification_df = gs_classification_df.sort_values(by='rank_test_score', ↳
    ascending=True)
gs_classification_df.head(20).round(4)
```

```
[ ]:   param_activation param_batch_size param_dropout_rate param_epochs \
2          relu          16          0.2          50
6          relu          16          0.3          50
10         relu          32          0.2          50
14         relu          32          0.3          50
18         relu          64          0.2          50
22         relu          64          0.3          50
4          relu          16          0.3          20
26         tanh          16          0.2          50
12         relu          32          0.3          20
42         tanh          64          0.2          50
0          relu          16          0.2          20
```

30	tanh	16	0.3	50
8	relu	32	0.2	20
46	tanh	64	0.3	50
38	tanh	32	0.3	50
20	relu	64	0.3	20
24	tanh	16	0.2	20
34	tanh	32	0.2	50
16	relu	64	0.2	20
36	tanh	32	0.3	20

	param_optimizer	mean_test_score	std_test_score	rank_test_score
2	adam	0.9068	0.0058	1
6	adam	0.9052	0.0054	2
10	adam	0.9038	0.0025	3
14	adam	0.9037	0.0033	4
18	adam	0.9031	0.0043	5
22	adam	0.9027	0.0043	6
4	adam	0.9009	0.0048	7
26	adam	0.9000	0.0052	8
12	adam	0.9000	0.0076	9
42	adam	0.8994	0.0056	10
0	adam	0.8994	0.0073	11
30	adam	0.8991	0.0050	12
8	adam	0.8988	0.0031	13
46	adam	0.8986	0.0064	14
38	adam	0.8985	0.0058	15
20	adam	0.8983	0.0065	16
24	adam	0.8982	0.0065	17
34	adam	0.8982	0.0054	18
16	adam	0.8974	0.0071	19
36	adam	0.8971	0.0061	20

2.8 Model evaluation

Compare the performance of the tuned model with the baseline model (i.e., the initial model without any hyperparameter tuning).

While a batch size of 32 and dropout of .3 yielded marginally better test score, + .0008, over 16 batch size, .2 dropout, it did so with over double the standard deviation in test score. For the tuned model, a dropout of .2 and batch size of 16 will be used for the greater consistency.

```
[ ]: activation = 'relu'
      batch_size = 16
      dropout_rate = 0.2
      epochs = 50
      optimizer = 'adam'
```

```
[ ]: tuned_classification_model = create_classification_model(activation=activation,
    ↪optimizer=optimizer, dropout_rate=dropout_rate)
```

```
[ ]: model_checkpoint = ModelCheckpoint('tuned_classification.h5',
    ↪save_best_only=True)

history = tuned_classification_model.fit(X_train_scaled, y_train,
    ↪epochs=epochs, batch_size=batch_size, validation_split=0.2, verbose=1,
    ↪callbacks=[early_stopping, model_checkpoint])
```

Epoch 1/50

```
c:\Users\Reed Oken\AppData\Local\Programs\Python\Python310\lib\site-
packages\keras\engine\data_adapter.py:1700: FutureWarning: The behavior of
`series[i:j]` with an integer-dtype index is deprecated. In a future version,
this will be treated as *label-based* indexing, consistent with e.g. `series[i]`
lookups. To retain the old behavior, use `series.iloc[i:j]`. To get the future
behavior, use `series.loc[i:j]`.
    return t[start:end]
```

```
826/826 [=====] - 2s 1ms/step - loss: 0.3517 -
accuracy: 0.8547 - val_loss: 0.2826 - val_accuracy: 0.8888
```

Epoch 2/50

```
826/826 [=====] - 1s 1ms/step - loss: 0.2975 -
accuracy: 0.8809 - val_loss: 0.2690 - val_accuracy: 0.8961
```

Epoch 3/50

```
826/826 [=====] - 1s 1ms/step - loss: 0.2854 -
accuracy: 0.8858 - val_loss: 0.2666 - val_accuracy: 0.8940
```

Epoch 4/50

```
826/826 [=====] - 1s 1ms/step - loss: 0.2788 -
accuracy: 0.8866 - val_loss: 0.2572 - val_accuracy: 0.8973
```

Epoch 5/50

```
826/826 [=====] - 1s 1ms/step - loss: 0.2706 -
accuracy: 0.8884 - val_loss: 0.2449 - val_accuracy: 0.9025
```

Epoch 6/50

```
826/826 [=====] - 1s 1ms/step - loss: 0.2645 -
accuracy: 0.8900 - val_loss: 0.2374 - val_accuracy: 0.9031
```

Epoch 7/50

```
826/826 [=====] - 1s 1ms/step - loss: 0.2577 -
accuracy: 0.8945 - val_loss: 0.2366 - val_accuracy: 0.9040
```

Epoch 8/50

```
826/826 [=====] - 1s 1ms/step - loss: 0.2547 -
accuracy: 0.8953 - val_loss: 0.2336 - val_accuracy: 0.9085
```

Epoch 9/50

```
826/826 [=====] - 1s 1ms/step - loss: 0.2535 -
accuracy: 0.8973 - val_loss: 0.2425 - val_accuracy: 0.8985
```

Epoch 10/50

```
826/826 [=====] - 1s 1ms/step - loss: 0.2534 -
accuracy: 0.8937 - val_loss: 0.2316 - val_accuracy: 0.9079
```

Epoch 11/50
826/826 [=====] - 1s 1ms/step - loss: 0.2510 -
accuracy: 0.8980 - val_loss: 0.2272 - val_accuracy: 0.9088

Epoch 12/50
826/826 [=====] - 1s 1ms/step - loss: 0.2494 -
accuracy: 0.8983 - val_loss: 0.2371 - val_accuracy: 0.9091

Epoch 13/50
826/826 [=====] - 1s 1ms/step - loss: 0.2468 -
accuracy: 0.8979 - val_loss: 0.2306 - val_accuracy: 0.9088

Epoch 14/50
826/826 [=====] - 1s 1ms/step - loss: 0.2485 -
accuracy: 0.8989 - val_loss: 0.2240 - val_accuracy: 0.9103

Epoch 15/50
826/826 [=====] - 1s 1ms/step - loss: 0.2444 -
accuracy: 0.8999 - val_loss: 0.2240 - val_accuracy: 0.9109

Epoch 16/50
826/826 [=====] - 1s 1ms/step - loss: 0.2431 -
accuracy: 0.9006 - val_loss: 0.2265 - val_accuracy: 0.9100

Epoch 17/50
826/826 [=====] - 1s 1ms/step - loss: 0.2441 -
accuracy: 0.8980 - val_loss: 0.2230 - val_accuracy: 0.9134

Epoch 18/50
826/826 [=====] - 1s 1ms/step - loss: 0.2402 -
accuracy: 0.9024 - val_loss: 0.2251 - val_accuracy: 0.9091

Epoch 19/50
826/826 [=====] - 1s 1ms/step - loss: 0.2430 -
accuracy: 0.9008 - val_loss: 0.2245 - val_accuracy: 0.9121

Epoch 20/50
826/826 [=====] - 1s 1ms/step - loss: 0.2400 -
accuracy: 0.8990 - val_loss: 0.2578 - val_accuracy: 0.8940

Epoch 21/50
826/826 [=====] - 1s 1ms/step - loss: 0.2392 -
accuracy: 0.8997 - val_loss: 0.2225 - val_accuracy: 0.9118

Epoch 22/50
826/826 [=====] - 1s 1ms/step - loss: 0.2378 -
accuracy: 0.9012 - val_loss: 0.2341 - val_accuracy: 0.9058

Epoch 23/50
826/826 [=====] - 1s 1ms/step - loss: 0.2382 -
accuracy: 0.9018 - val_loss: 0.2495 - val_accuracy: 0.9018

Epoch 24/50
826/826 [=====] - 1s 1ms/step - loss: 0.2336 -
accuracy: 0.9044 - val_loss: 0.2172 - val_accuracy: 0.9143

Epoch 25/50
826/826 [=====] - 1s 1ms/step - loss: 0.2340 -
accuracy: 0.9022 - val_loss: 0.2197 - val_accuracy: 0.9170

Epoch 26/50
826/826 [=====] - 1s 1ms/step - loss: 0.2366 -
accuracy: 0.9028 - val_loss: 0.2196 - val_accuracy: 0.9131


Epoch 27/50
826/826 [=====] - 1s 1ms/step - loss: 0.2321 - accuracy: 0.9009 - val_loss: 0.2140 - val_accuracy: 0.9173
Epoch 28/50
826/826 [=====] - 1s 1ms/step - loss: 0.2315 - accuracy: 0.9041 - val_loss: 0.2147 - val_accuracy: 0.9176
Epoch 29/50
826/826 [=====] - 1s 1ms/step - loss: 0.2292 - accuracy: 0.9044 - val_loss: 0.2166 - val_accuracy: 0.9173
Epoch 30/50
826/826 [=====] - 1s 1ms/step - loss: 0.2314 - accuracy: 0.9050 - val_loss: 0.2137 - val_accuracy: 0.9140
Epoch 31/50
826/826 [=====] - 1s 1ms/step - loss: 0.2286 - accuracy: 0.9049 - val_loss: 0.2390 - val_accuracy: 0.9006
Epoch 32/50
826/826 [=====] - 1s 1ms/step - loss: 0.2273 - accuracy: 0.9025 - val_loss: 0.2150 - val_accuracy: 0.9173
Epoch 33/50
826/826 [=====] - 1s 1ms/step - loss: 0.2258 - accuracy: 0.9048 - val_loss: 0.2594 - val_accuracy: 0.8967
Epoch 34/50
826/826 [=====] - 1s 1ms/step - loss: 0.2256 - accuracy: 0.9062 - val_loss: 0.2133 - val_accuracy: 0.9182
Epoch 35/50
826/826 [=====] - 1s 1ms/step - loss: 0.2259 - accuracy: 0.9057 - val_loss: 0.2256 - val_accuracy: 0.9134
Epoch 36/50
826/826 [=====] - 1s 1ms/step - loss: 0.2220 - accuracy: 0.9079 - val_loss: 0.2146 - val_accuracy: 0.9152
Epoch 37/50
826/826 [=====] - 1s 1ms/step - loss: 0.2224 - accuracy: 0.9093 - val_loss: 0.2114 - val_accuracy: 0.9197
Epoch 38/50
826/826 [=====] - 1s 1ms/step - loss: 0.2242 - accuracy: 0.9068 - val_loss: 0.2178 - val_accuracy: 0.9143
Epoch 39/50
826/826 [=====] - 1s 1ms/step - loss: 0.2213 - accuracy: 0.9062 - val_loss: 0.2135 - val_accuracy: 0.9191
Epoch 40/50
826/826 [=====] - 1s 1ms/step - loss: 0.2233 - accuracy: 0.9088 - val_loss: 0.2154 - val_accuracy: 0.9164
Epoch 41/50
826/826 [=====] - 1s 1ms/step - loss: 0.2185 - accuracy: 0.9092 - val_loss: 0.2171 - val_accuracy: 0.9155
Epoch 42/50
826/826 [=====] - 1s 1ms/step - loss: 0.2204 - accuracy: 0.9073 - val_loss: 0.2100 - val_accuracy: 0.9200


```


Epoch 43/50
826/826 [=====] - 1s 1ms/step - loss: 0.2206 -
accuracy: 0.9083 - val_loss: 0.2176 - val_accuracy: 0.9206
Epoch 44/50
826/826 [=====] - 1s 1ms/step - loss: 0.2207 -
accuracy: 0.9092 - val_loss: 0.2150 - val_accuracy: 0.9146
Epoch 45/50
826/826 [=====] - 1s 1ms/step - loss: 0.2187 -
accuracy: 0.9089 - val_loss: 0.2143 - val_accuracy: 0.9182
Epoch 46/50
826/826 [=====] - 1s 1ms/step - loss: 0.2188 -
accuracy: 0.9093 - val_loss: 0.2136 - val_accuracy: 0.9215
Epoch 47/50
826/826 [=====] - 1s 1ms/step - loss: 0.2186 -
accuracy: 0.9080 - val_loss: 0.2146 - val_accuracy: 0.9176

```

```
[ ]: training_plots(history, 'accuracy')
```



```
[ ]: training_plots(history, 'loss')
```



```
[ ]: loss, accuracy = tuned_classification_model.evaluate(X_test_scaled, y_test, verbose=0)
print(f'Test Loss: {loss:.4f}')
print(f'Test accuracy: {accuracy:.4f}')
```

Test Loss: 0.2053
Test accuracy: 0.9193

```
[ ]: predictions = tuned_classification_model.predict(X_test_scaled)
threshold = 0.5
y_pred = np.where(predictions >= threshold, 1, 0)
```

129/129 [=====] - 0s 600us/step

2.8.1 Threshold tuning

```
[ ]: thresh_arr = []
prec_arr = []
recall_arr = []
f1_arr = []
```

```

acc_arr = []

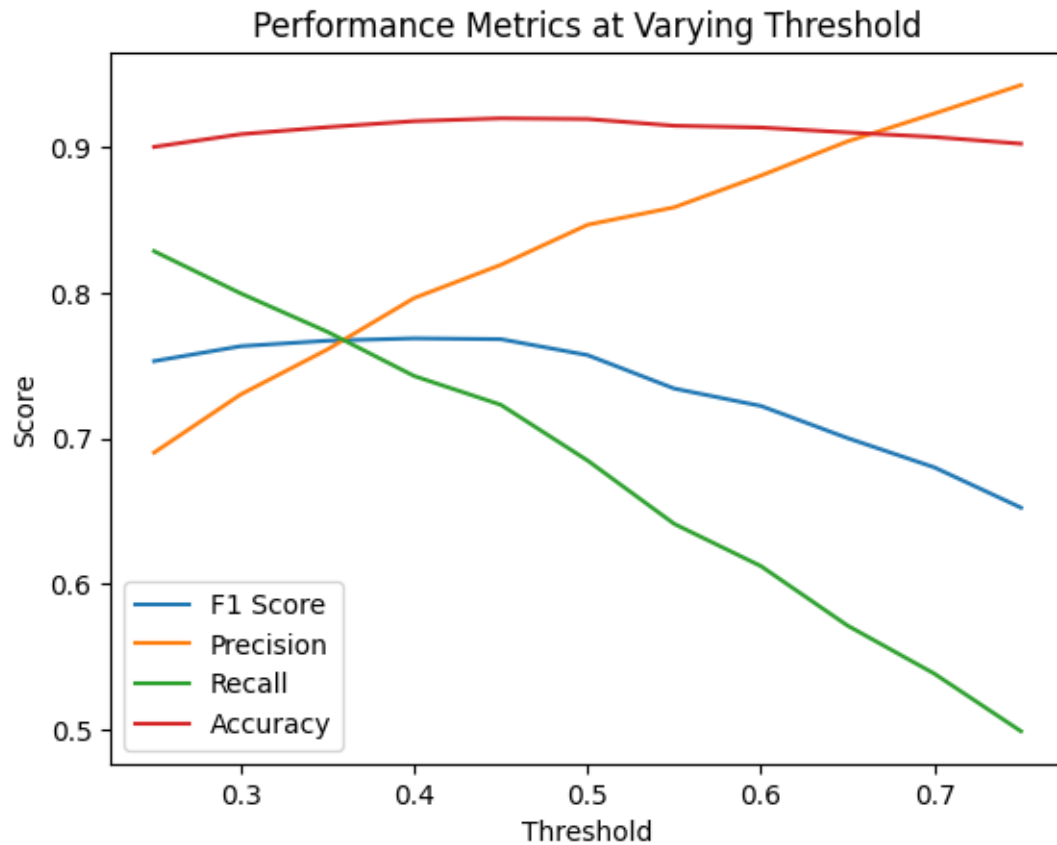
for i in range(11):
    thresh = .25 + i*.05
    thresh_arr.append(thresh)
    y_pred = np.where(predictions >= thresh, 1, 0)
    precision = precision_score(y_test, y_pred)
    recall = recall_score(y_test, y_pred)
    f1 = f1_score(y_test, y_pred)
    acc = accuracy_score(y_test, y_pred)
    prec_arr.append(precision)
    recall_arr.append(recall)
    f1_arr.append(f1)
    acc_arr.append(acc)

```

```

[ ]: sns.lineplot(x=thresh_arr, y=f1_arr, label='F1 Score')
sns.lineplot(x=thresh_arr, y=prec_arr, label='Precision')
sns.lineplot(x=thresh_arr, y=recall_arr, label='Recall')
sns.lineplot(x=thresh_arr, y=acc_arr, label='Accuracy')
plt.legend()
plt.xlabel('Threshold')
plt.ylabel('Score')
plt.title('Performance Metrics at Varying Threshold')
plt.show()


```



Threshold of .5 appears to provide the best performance across all performance metrics

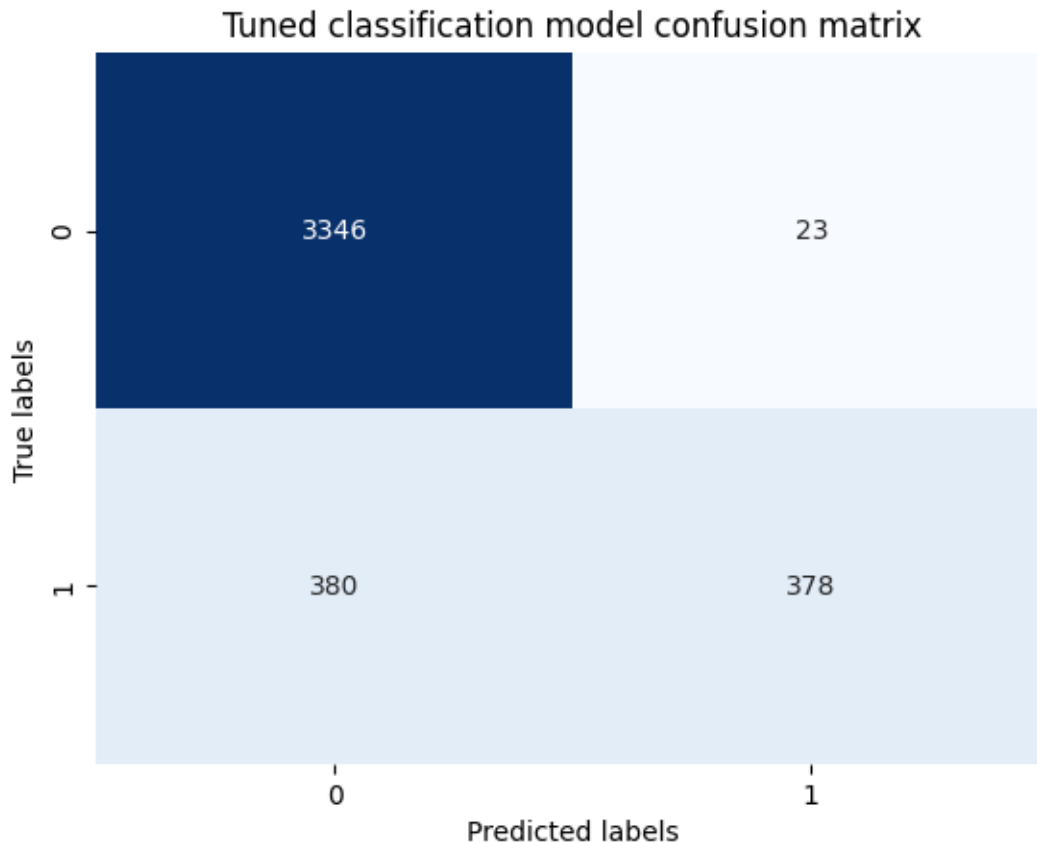
2.8.2 Test performance on tuned model

```
[ ]: sns.histplot(predictions)
plt.title('Tuned Classification Model Probability Histogram')
plt.xlabel('Probability')
plt.show()
```



```
[ ]: confusion_mat = confusion_matrix(y_test, y_pred)

sns.heatmap(confusion_mat, annot=True, cmap='Blues', fmt='g', cbar=False)
plt.xlabel('Predicted labels')
plt.ylabel('True labels')
plt.title('Tuned classification model confusion matrix')
plt.show()
```



```
[ ]: print(f'Baseline classification model test accuracy: {baseline_acc:.4f}')
      print(f'Tuned classification model test accuracy: {tuned_acc:.4f}\n')

      print(f'Baseline classification model test precision: {baseline_precision:.4f}')
      print(f'Tuned classification model test precision: {tuned_precision:.4f}\n')

      print(f'Baseline classification model test recall: {baseline_recall:.4f}')
      print(f'Tuned classification model test recall: {tuned_recall:.4f}\n')

      print(f'Baseline classification model test F1: {baseline_f1:.4f}')
      print(f'Tuned classification model test F1: {tuned_f1:.4f}')
```

Baseline classification model test accuracy: 0.9159
Tuned classification model test accuracy: 0.9198

Baseline classification model test precision: 0.8236
Tuned classification model test precision: 0.8280

Baseline classification model test recall: 0.6900
Tuned classification model test recall: 0.7111

Baseline classification model test F1: 0.7509
Tuned classification model test F1: 0.7651